

Prometheus Cheatsheet

TJ ROBINSON

November 29, 2025

1 METRIC VS. TIME-SERIES

Metric is a numerical value that represents a specific measurement at a specific point in time.

Time Series is a sequence of data points, typically ordered by time, that represents the evolution of a metric over time.

- every metric & label combo is a time series.

2 PROMETHEUS TSDB

Prometheus TSDB is a time series database designed for storing and querying time series data, particularly metrics data collected from various sources.

2.1 HEAD BLOCK

Head block is the in-memory storage component of Prometheus TSDB that holds the most recent time series data before it is flushed to disk.

- Each chunk for a time series in the head block typically contains data points for a specific time range and can grow to 120 samples before it is flushed to the active chunk's segment file in the `chunks_head` directory on the disk.
- Each head chunk segment file can grow to 128 MB before it is finalized and a new segment file is created for that chunk.
- When the head block reaches 1.5x the minimum block size (default 2 hours), a new block is created, and the current head block is flushed to disk as a finalized block.

2.2 WAL

WAL (Write-Ahead Log) is a log file that records all changes made to the head block before they are applied, ensuring data durability and consistency in case of crashes.

- The WAL prevents loss of recent data by allowing Prometheus to replay the log and restore the head block to its last known state upon restart.
- WAL segments are stored in the `wal` directory and are named sequentially (e.g., 00000000, 00000001, etc.).
- Each WAL segment can grow up to 512 MB before a new segment is created

3 PROMQL (PROMETHEUS QUERY LANGUAGE)

promQL is the query language used to query metrics stored in Prometheus.

PromQL queries can evaluate to one of four types of results:

- Instant vector: A set of time series containing a single sample for each time series, all sharing the same timestamp.
- Range vector: A set of time series containing a range of data points over time for each time series.
- Scalar: A single numerical value. (i.e., 177.6)
- String: A single string value. (i.e., "Primus Sux")

3.1 SELECTORS & MATYCHERS

3.1.1 Selctors

- `timeseries`
- Range Vector Selector will returns all values for a metric over a time period

3.1.2 Label Matchers

- Label matchers allow you to return metrics for a subset of time series.
 - = Exact match on a label value.
 - != Negative equality matcher - return time series that do not have specified label.
 - =~ Regular Expression Matcher - utilize `regex` to match time series.
 - !~ Negative Regular Expression Matcher - return time series that do not match `regex` expression
- multiple seletors can. be used in a query by seperating them with a comma.

3.2 MODIFIERS

offset modifiers allows you top retrieve historc data.

@ modifier allows you to retrieve data from a specific point in time.

- **offset** and **@** modifiers can be combined in a query (order does not matter).
- **Range vectors** can also be combined with **offset** and **@** modifiers.